


# The Power of Image Filtering...

## Hybrid Imagery Example



Hybrid Images, Oliva et al., [http://cvcl.mit.edu/publications/OlivaTorralba\\_Hybrid\\_Siggraph06.pdf](http://cvcl.mit.edu/publications/OlivaTorralba_Hybrid_Siggraph06.pdf)





What is an image?






# What is an image?



**Digital Camera**




**The Eye**

**We'll focus on these in this class**

**(More on this process later)**

# What is an image?

- A grid (matrix) of intensity values



=

255	255	255	255	255	255	255	255	255	255	255	255
255	255	255	255	255	255	255	255	255	255	255	255
255	255	255	20	0	255	255	255	255	255	255	255
255	255	255	75	75	75	255	255	255	255	255	255
255	255	75	95	95	75	255	255	255	255	255	255
255	255	96	127	145	175	255	255	255	255	255	255
255	255	127	145	175	175	175	255	255	255	255	255
255	255	127	145	200	200	175	175	95	255	255	255
255	255	127	145	200	200	175	175	95	47	255	255
255	255	127	145	145	175	127	127	95	47	255	255
255	255	74	127	127	127	95	95	95	47	255	255
255	255	255	74	74	74	74	74	74	255	255	255
255	255	255	255	255	255	255	255	255	255	255	255
255	255	255	255	255	255	255	255	255	255	255	255


(common to use one byte per value: 0 = black, 255 = white)



# What is an image?

---

- We can think of a (grayscale) image as a **function**,  $f$ , from  $\mathbb{R}^2$  to  $\mathbb{R}$  (or a 2D *signal*):
  - $f(x,y)$  gives the **intensity** at position  $(x,y)$



- A **digital** image is a discrete (**sampled, quantized**) version of this function



# Image transformations

---

- As with any function, we can apply operators to an image



$$g(x,y) = f(x,y) + 20$$



$$g(x,y) = f(-x,y)$$

- We'll talk about a special kind of operator, *convolution* (linear filtering)



# Question: Noise reduction

- Given a camera and a still scene, how can you reduce noise?



Take lots of images and average them!

What's the next best thing?



# Image filtering

---

- Modify the pixels in an image based on some function of a local neighborhood of each pixel

10	5	3
4	5	1
1	1	7

Local image data

Some function




	7	

Modified image data



# Linear filtering

- One simple version: linear filtering (cross-correlation, convolution)
  - Replace each pixel by a linear combination of its neighbors
- The prescription for the linear combination is called the “kernel” (or “mask”, “filter”)





# Convolution

---

- Same as cross-correlation, except that the kernel is “flipped” (horizontally and vertically)

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i - u, j - v]$$

- Convolution is **commutative** and **associative**

This is called a **convolution** operation:


$$G = H * F$$



# Kernel Transformations

Replace each pixel by a linear combination of its neighbors

- The prescription for the linear combination is called the "kernel" (or "mask", "filter")
- Involves sliding a kernel (filter) across an image.
- A mask should always be in odd number, because otherwise you cannot find the mid of the mask.





## How to do it

---

In order to perform a filter on an image, following steps should be taken.

- 1) Slide the mask onto the image.
- 2) Multiply the corresponding elements and then add them
- 3) Repeat this procedure until all values of the image have been calculated.



# Linear filters: examples


---



Original



0	0	0
0	1	0
0	0	0



Identical image



# Linear filters: examples

---



Original



0	0	0
1	0	0
0	0	0



Shifted left  
By 1 pixel



# Linear filters: examples

---




Original



$\frac{1}{9}$

1	1	1
1	1	1
1	1	1



Blur (with a mean filter)



2D Average filtering example using a 3 x 3 sampling window:

Keeping border values unchanged

$$\text{Average} = \text{round}(1+4+0+2+2+4+1+0+1)/9 = 2$$

Input

1	4	0	1	3	1
2	2	4	2	2	3
1	0	1	0	1	0
1	2	1	0	2	2
2	5	3	1	2	5
1	1	4	2	3	0

Output

1	4	0	1	3	1
2	2	2	2	1	3
1	2	1	1	1	0
1	2	1	1	1	2
2	2	2	2	2	5
1	1	4	2	3	0



---

We just did blurring with a 3x3 mask.

Consider a 5x5 mask:

- 1) What would the effect on blurring be?
- 2) What would the values of the mask need to be?

It is also possible to have masks that are 7x7, 9x9, etc... Try it out!



## Octave Implementation

```
rgbmonkey = imread("monkey.jpeg");  
  
gs = mat2gray(rgb2gray(rgbmonkey)); %mat2gray for normalisation  
  
kernel_identity = [0 0 0;0 1 0;0 0 0];  
  
resultant = filter2 (kernel_identity,gs);
```

Try implementing different levels of blurring!



## Octave built-in blur function

---

*resultant = imsmooth(img, type, options)*

Where type can be:

"Gaussian" - This is the default.



"Average" - Smoothing using a rectangular averaging (mean) filter.

"Disk" - Smoothing using a circular averaging (mean) filter.

....



# Gaussian Kernel




$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

# Gaussian filters



---




 = 1 pixel



 = 5 pixels



 = 10 pixels




 = 30 pixels




## High-pass filters: Sharpening

---



before



after



# Sharpening

- What does blurring take away?



Let's add it back:



Source: S. Lazebnik



# Linear filters: examples



Original

$$* \left( \begin{pmatrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{pmatrix} - \frac{1}{9} \begin{pmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{pmatrix} \right) =$$



**Sharpening filter**  
(accentuates edges)



# Sharpening revisited

- What does blurring take away?



—



=



Let's add it back:



+  $\alpha$




=






# Sharpen filter



scaled impulse


$$F + \alpha (F - F * H)$$

image                      blurred image



Gaussian

$\approx$



# Sharpen filter

---

unfiltered

filtered





# Convolution in the real world

## Camera shake



Source: Fergus, *et al.* "Removing Camera Shake from a Single Photograph", SIGGRAPH 2006

**Bokeh:** Blur in out-of-focus regions of an image.



Source: <http://lullaby.homepage.dk/diy-camera/bokeh.html>